# Drop, Reflect and Reuse a Quantum Temporary

ANONYMOUS AUTHOR(S)

We characterize safe drops by contextual equivalence. Without measurement, a safe drop means the temporary ends in the same pure state at every input, so at most one reflection restores it. The criterion covers records and loans, providing the reuse check in our OpenQASM 3 compiler.

## 1 Specifying a safe drop

When a temporary qubit is no longer needed, we want to reuse it without disturbing the rest of the computation. The standard technique, uncomputation, runs the computation that produced a temporary backwards, which can take as many gates as the computation itself [14, §1]. Whether a temporary may be safely discarded is a separate question, which languages and tools such as Silq and Unqomp settle with conservative rules for automatic reclamation that can reject safe temporaries [2, §4.4], [14, §6.2].

Consider $p$ below, which allocates its only temporary `t`, of dimension $d$, in pure state $\psi$, runs a measurement-free computation $M$ on `(t,v)`, and drops `t`. We compare $p$ just before the drop with an independent witness $\mathcal{S}(P,U)$, a comparison program that uses $P$ to prepare `t` without using the input `v`, and applies $U$ separately to `v`:

$$p(\texttt{v}) = \texttt{new t : [d] @ } \psi \texttt{; M; drop t; return v}$$

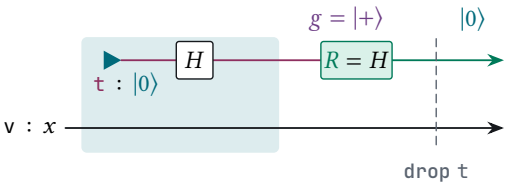
$$\mathcal{S}(P,U)(\texttt{v}) = \texttt{new t : [d] @ } \psi \texttt{; P t; U v; drop t; return v}$$

We call the drop *safe* if $p \equiv \mathcal{S}(P,U)$ for one witness on all inputs, meaning that no context can distinguish the programs stopped before their drops, with `t` and `v` returned.

We show that at a safe drop, `t` ends in one fixed pure state $g$ at every input. From $g$ we construct a return gate $R$ that restores $\psi$ on `t` alone with at most one reflection. The two examples below take a qubit `v`, with inserted return gates shown in green. To check each drop, we use the shaded *exposed prefix* before the chosen `drop` *site* (⫞). Its map $K$ sends input $x$ to the joint state there, with `t` first. A wire label `v : x` gives the register and its input state. Triangles (▶) mark allocation and ground symbols (⏚) end scope.

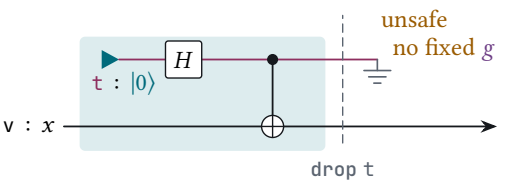
```
prog a (v : [2]) : [2] {
  new t : [2] @ |0>; H t;
  drop t; return v }
```

$K^{(a)}x = |+\rangle \otimes x$




```
prog b (v : [2]) : [2] {
  new t : [2] @ |0>; H t; CX t v;
  drop t; return v }
```

$K^{(b)}x = (|0\rangle \otimes x + |1\rangle \otimes Xx)/\sqrt{2}$




In our language, `drop` ends `t`'s scope without changing the joint state. Program (a) leaves `t` in $|+\rangle$ at every input, so $H$ alone restores `t` to $|0\rangle$. In program (b), `CX t v` flips `v` when `t` is $|1\rangle$, giving the Bell state $(|00\rangle + |11\rangle)/\sqrt{2}$ when `v` starts in $|0\rangle$. Both temporaries have the same lexical scope, so checking scopes would accept both drops. Silq and Unqomp's

`qfree` rules exclude $H$ because it takes basis states to superpositions, and reject both drops. Unqomp also rejects `H t; H t`, where `t` ends in $|0\rangle$ [14, §6.2].

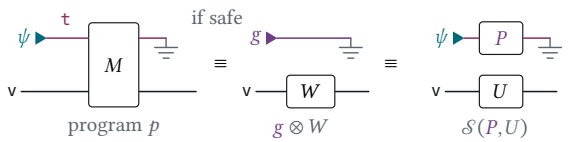
## 2 Checking the equivalence

$$K^{(a)} = \frac{1}{\sqrt{2}} \overbrace{\begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & -1 & 0 \\ 0 & 1 & 0 & -1 \end{bmatrix}}^{H \otimes I} \cdot \overbrace{\begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}}^{|0\rangle \otimes I} = \overbrace{\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 0 & 1 \end{bmatrix}}^{K^{(a)}} \begin{matrix} K_0 \\ K_1 \end{matrix} \quad \text{slices } I/\sqrt{2},\ I/\sqrt{2}$$

$$\overbrace{\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 0 & 1 \\ 1 & 0 \end{bmatrix}}^{K^{(b)} = \mathrm{CX}\,K^{(a)}} \begin{matrix} K_0 \\ K_1 \end{matrix} \quad \text{slices } I/\sqrt{2},\ X/\sqrt{2}$$

To read the joint state from the temporary's side before the drop, we group the exposed prefix's matrix rows by `t`'s basis value. For program (a), allocation $|0\rangle \otimes I$ acts before $H \otimes I$. The blue *slice* $K_0$ gives $|0\rangle \otimes K_0 x$ and the gold $K_1$ gives $|1\rangle \otimes K_1 x$. The joint state is the sum of these two terms.

A *context* is a closed program whose hole takes `v` to `(t,v)`. It supplies an input, possibly entangled with a reference, applies gates and reads all registers, recording outcome weights, as in




```
new v : [2] @ |0>; (t,v) = { □ }; CX t v; H t; return (t,v).
```

This context undoes program (b)'s gates and reads 00 with weight 1. An independent witness leaves `t` and `v` uncorrelated, so the same context reads 00 at most half the time. The witness prepares `t` in a fixed pure state $g$ without using the input, while $U$ computes a map $W$ on the rest, sending input $x$ to their joint state $Wx$. The witness therefore has slice $g_i W$ at basis value $i$, where $g_i$ is the coefficient of $|i\rangle$ in $g$. For $p$ to be equivalent to the witness, no context can distinguish their joint states before the drop, on any input. This requires the same slice form $K_i = g_i W$ for $p$, giving $Kx = g \otimes Wx$. In program (a), equal slices give $g = |+\rangle$ and $W = I$, while in program (b) no scalar turns $I$ into $X$.

The witness prepares the dropped registers together as one *block*, without the input, and allocates matching registers for $p$'s other temporaries so contexts can compare the full joint state, including hidden registers.

Theorem 2.1 (without measurement). *Let $K$ be the exposed map before dropping allocated temporaries with pure allocation state $\psi$ in a well-typed, measurement-free program $p$. These are equivalent.*

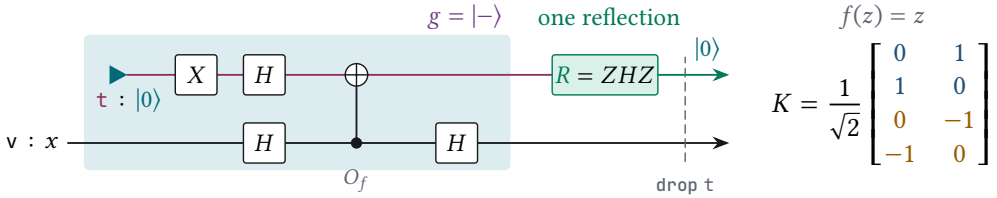
*(1) $p \equiv \mathcal{S}(P,U)$ for some independent witness.*

*(2) $K = g \otimes W$ for a unit vector $g$ on the block and an isometry $W$, a length-preserving map from the input to all other registers.*

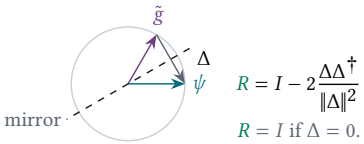
*(3) One unitary on the block alone restores $\psi$ at every input.*

Deutsch–Jozsa shows why slices need only be proportional. Its oracle $O_f$ flips `t` when $f(z) = 1$. In the minus state $|-\rangle = (|0\rangle - |1\rangle)/\sqrt{2}$, a flip only negates `t`, so the oracle leaves

it in $|-\rangle$ and gives the data state $|z\rangle$ a sign $(-1)^{f(z)}$. These signs define $D_f|z\rangle = (-1)^{f(z)}|z\rangle$. For one data qubit, the Hadamards give $W = HD_fH$, pictured for $f(z) = z$. The slices are $W/\sqrt{2}$ and $-W/\sqrt{2}$, proportional for every $f$, with $g = |-\rangle$. Reversing the preparation takes $H$ then $X$ on `t`. The reflection $R = ZHZ$ restores `t` in one gate, and in this case both methods leave the data unchanged.




To build $R$, choose a phase representative $\tilde{g}$ of our *certificate* $g$ with real, nonnegative overlap with $\psi$, moving the phase into $W$. The difference $\Delta = \psi - \tilde{g}$ points from $\tilde{g}$ to $\psi$. If they coincide, no gate is needed. Otherwise the Householder reflection through the mirror orthogonal to $\Delta$ sends $\tilde{g}$ to $\psi$ [8]. The compiler aligns $\Delta/\|\Delta\|$ with the usual allocation state $|0\cdots0\rangle$, flips only that state's sign and undoes the alignment. Its QASM circuit may use several gates, but depends on the block's certificate rather than the producer's length. If $\Delta/\|\Delta\|$ is the uniform superposition of the block's basis states, then $-R$ is Grover's diffusion.



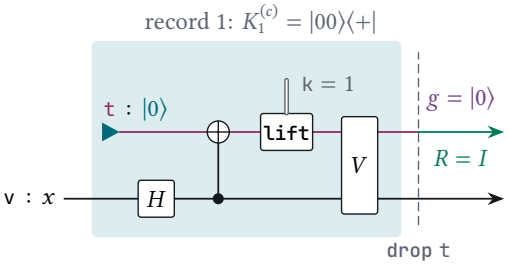

## 3 One measurement record at a time

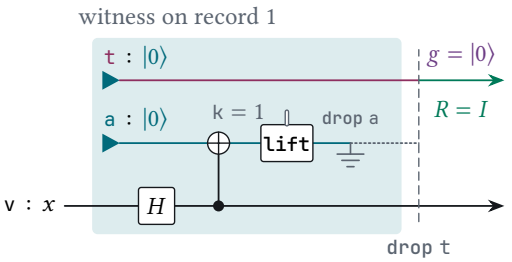
With measurement, we check each outcome sequence, or *record* $o$, separately. Our `lift[S] t` keeps `t` live and binds 1 for membership in a basis set $S$, or 0 otherwise. Only classical `if` guards read this outcome, as in dynamic lifting [6]. Thus `lift[{0}] t` returns 1 when `t` is read as 0. The exposed map $K_o$ stays unnormalized, so $\|K_o x\|^2$ is the record's probability at a normalized input $x$.

Program (c) uses a two-qubit unitary $V$ that fixes $|00\rangle$ and sends $|11\rangle$ to $(|01\rangle + |10\rangle)/\sqrt{2}$. Its independent witness measures a fresh helper qubit `a` instead of `t`:

```
prog c (v : [2]) : [2] {
  new t : [2] @ |0>; H v; CX v t;
  k = lift[{0}] t; V t v;
  drop t; return v }
```

```
prog witness (v : [2]) : [2] {
  new t : [2] @ |0>; new a : [2] @ |0>;
  H v; CX v a; k = lift[{0}] a;
  drop a; drop t; return v }
```



At input $x = \alpha|0\rangle + \beta|1\rangle$, measurement followed by $V$ leaves these unnormalized states:

- Record 1 (safe) leaves $(\alpha + \beta)|00\rangle/\sqrt{2}$, with weight $|\alpha + \beta|^2/2$. The temporary remains in $|0\rangle$, so $R = I$.
- Record 0 (unsafe) leaves $(\alpha - \beta)(|01\rangle + |10\rangle)/2$, with weight $|\alpha - \beta|^2/2$. Its nonzero outputs normalize to a Bell pair, so this record has no pure certificate $g$.

On record 1, $K_1^{(c)} = |0\rangle\otimes W_1$ with $W_1 = |0\rangle\langle+|$. The map $W_1$ keeps only the $|+\rangle$ component of the input, so it never lengthens a vector and is a *contraction*. For example, $W_1$ gives $|0\rangle$ with weight 1 at $x = |+\rangle$, but $|0\rangle/\sqrt{2}$ with weight $1/2$ at $x = |0\rangle$. Normalizing would give $|0\rangle$ in both cases, hiding the dependence of record 1's weight on the input.

If we allocate the witness's `t` and `a` in $|0\rangle$ and apply `H v; CX v a`, then these steps preserve total weight, so reproducing record 1 requires measurement. They leave `t` in $|0\rangle$, while `v` and `a` have the same value in each basis term. Measuring the independent `t` would give the same probability at every input, so the witness measures `a`. Outcome 1 leaves `t` and `a` in $|0\rangle$ and `v` in $W_1 x$, with record 1's weight. Program (c)'s unsafe record 0 still prevents reuse.

The same construction realizes any contraction $W_o$ by giving the helper's other outcome the missing weight $\|x\|^2 - \|W_o x\|^2$. With fresh registers matching input and output spaces, the two outcomes preserve norm and extend to a unitary $U$ [17, Thm. 2]. One `lift` then selects $W_o x$, while $P$ prepares $g$ from $\psi$ on the block.

We call record $o$ *safe* and write $p \equiv_o \mathcal{S}(P,U)$ if no context distinguishes $K_o$ from the witness's map $g \otimes W_o$ on kept outcome 1 before the drop, omitting only its fixed $|0\rangle$ helper. Both maps stay unnormalized to retain record probabilities, whatever the other records do.

Theorem 3.1 (on a selected record). *Let $K_o$ be the exposed map before dropping allocated temporaries on a record $o$ of a well-typed program $p$. These are equivalent.*

*(1) $p \equiv_o \mathcal{S}(P,U)$ for some independent witness using one* `lift`.
*(2) $K_o = g \otimes W_o$ for a unit vector $g$ on the block and a contraction $W_o$ to all other registers.*
*(3) One unitary on the block alone restores its pure allocation state $\psi$ at every input on $o$.*

*If $K_o = 0$, then the record has probability zero at every input and all three conditions hold. Return gates may differ between records.*

Since reuse must preserve observable behavior, we might hope to decide safety from the *summed channel*, the mixture of returned-register and reference states over records [20, §2.3.2]. With `t` prepared in $|+\rangle$, applying controlled $Z$ to `(t,v)` before measuring `t` in the standard basis gives a safe drop on both records. Measuring `t` first, then flipping `t` under `v`'s control, gives an unsafe drop on both records. Both drop `t` and return `v`, giving the same unread standard-basis measurement with weight $1/2$ for each record at every input.

## 4 Blocks and loans

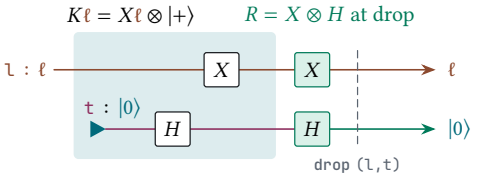
An allocated block in a fixed pure state is safe to drop even when its parts are mixed, as with a fresh Bell pair. We also consider blocks containing *loans*, registers borrowed in unknown states, possibly entangled with their caller, who expects them back unchanged.

The witness $\mathcal{S}(P,U)$ receives the unknown loan state $\ell$, so its block factor is an isometry $B$ rather than one fixed state $g$. A drop that names loans is safe when $K = B \otimes W$, with $B$ acting on the loans and $U$ computing $W$ on the other inputs $x$.




With measurement, we again check each record $o$ for $K_o = B_o \otimes W_o$, allowing $R$ to depend on $o$. Without measurement, $R$ restores the loan's normalized joint state with its caller reference, and the same holds on each occurring record when loan and reference are uncorrelated with other inputs. Otherwise measuring those inputs can change that state. The diagram below gives $K$ and $R$ for `lend`, whose unit return is `() : 1`.

```
prog lend (loan l : [2]) : 1 {
  new t : [2] @ |0>;
  H t; X l;
  drop (l,t); return () }
```




## 5 Scope and further work

We prove progress and preservation for a linear, first-order abstract machine with lifting and classical guards. For reuse, it restores safe blocks and pools temporaries for `new` of matching dimension and state, preserving records, weights and the reduced state of returned registers on each record. To decide safety, our exact Haskell checker tests each record's basis inputs for a common temporary state or loan isometry $B$, taking time exponential in the number of qubits. With only basis allocations and $X$, `CX` and Toffoli, safety of drops without loans is coNP-complete, even without measurement and with one drop (compare [11, Cor. 3.1]).

Our OpenQASM 3 compiler runs this check, computes numerical return gates, and reuses blocks safe on every record. Only basis allocations compile, program (c)'s $V$ is unsupported, and no loan may be dropped after a `lift` or an `if`. For the general loan returns it supports, the compiler checks $RB = J$ numerically, with entrywise tolerance $10^{-9}$. Measuring and resetting a qubit to $|0\rangle$ also permits reuse [4, 9], but destroys an unknown loan state.

Where Qurts and Reqomp reclaim by uncomputing [7, 15], our test constructs a return from the state at the drop, certifying the balanced Deutsch–Jozsa temporary that Qutes does not reclaim [5, Defs. 2–6]. Twist verifies annotated purity through static analysis and runtime separability tests [22, §1.3], but one return gate needs a pure state common to all inputs. The gate built from that state provides the restoration that ReQWIRE verifies [16, §4] and that Q# and fractional types require [3, 13].

For loans, prior work repairs separable errors on a dirty ancilla by local rotations [10, §5.2, Alg. 2], while our gate undoes the block's isometry $B$ at the drop. QBorrow also requires unchanged loans after each denoted operation [18, Def. 5.1] and Bona after each borrowing body [21, Def. 3.2], a handback property that separation logics verify deductively [12, 19].

Our language excludes loops and quantum recursion, so we leave open when a temporary used in a loop or recursive call is safe to drop. We also allow all unitaries as gates, although a fixed set may lack the return. For example, circuits over $X$, `CX`, Toffoli and $H \otimes H$ have dyadic entries [1, Cor. 5.6]. Applying $H \otimes H$ to `(t,v)` with `t` allocated in $|0\rangle$ gives a safe drop, yet restoring `t` alone requires non-dyadic entries $\pm 1/\sqrt{2}$.

## References


[1] Matthew Amy, Andrew N. Glaudell, and Neil J. Ross. 2020. Number-Theoretic Characterizations of Some Restricted Clifford+T Circuits. *Quantum* 4, Article 252 (2020). doi:10.22331/q-2020-04-06-252

[2] Benjamin Bichsel, Maximilian Baader, Timon Gehr, and Martin Vechev. 2020. Silq: A High-Level Quantum Language with Safe Uncomputation and Intuitive Semantics. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '20)*. ACM, 286–300. doi:10.1145/3385412.3386007

[3] Chao-Hong Chen, Vikraman Choudhury, Jacques Carette, and Amr Sabry. 2020. Fractional Types: Expressive and Safe Space Management for Ancilla Bits. In *Reversible Computation — 12th International Conference, RC 2020 (Lecture Notes in Computer Science, Vol. 12227)*. Springer, 169–186. doi:10.1007/978-3-030-52482-1_10

[4] Matthew DeCross, Eli Chertkov, Megan Kohagen, and Michael Foss-Feig. 2023. Qubit-Reuse Compilation with Mid-Circuit Measurement and Reset. *Physical Review X* 13, 4 (2023), 041057. doi:10.1103/PhysRevX.13.041057

[5] Simone Faro, Francesco Pio Marino, and Gabriele Messina. 2026. Reversible Lifetime Semantics for Quantum Programs. arXiv:2603.14538v1 [quant-ph]

[6] Peng Fu, Kohei Kishida, Neil J. Ross, and Peter Selinger. 2023. Proto-Quipper with Dynamic Lifting. *Proceedings of the ACM on Programming Languages* 7, POPL, Article 11 (2023), 26 pages. doi:10.1145/3571204

[7] Kengo Hirata and Chris Heunen. 2025. Qurts: Automatic Quantum Uncomputation by Affine Types with Lifetime. *Proceedings of the ACM on Programming Languages* 9, POPL, Article 6 (2025), 28 pages. doi:10.1145/3704842

[8] Alston S. Householder. 1958. Unitary Triangularization of a Nonsymmetric Matrix. *J. ACM* 5, 4 (Oct. 1958), 339–342. doi:10.1145/320941.320947

[9] Fei Hua, Yuwei Jin, Yanhao Chen, Suhas Vittal, Kevin Krsulich, Lev S. Bishop, John Lapeyre, Ali Javadi-Abhari, and Eddy Z. Zhang. 2023. CaQR: A Compiler-Assisted Approach for Qubit Reuse through Dynamic Circuit. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3 (ASPLOS '23)*. 59–71. doi:10.1145/3582016.3582030

[10] Jiqi Li, Jingyi Mei, Wang Fang, and Ji Guan. 2026. Formal Verification of Quantum Ancilla Safety. In *Computer Aided Verification – CAV 2026, Part III (Lecture Notes in Computer Science, Vol. 16684)*. Springer Nature Switzerland, 326–348. doi:10.1007/978-3-032-32537-2_16

[11] Chenke Liu, Li Zhou, and Boning Meng. 2026. Quantum Uncomputation of Clean and Dirty Ancilla Qubits. *Proceedings of the ACM on Programming Languages* 10, OOPSLA2 (2026), 2270–2298. doi:10.1145/3839528

[12] Yusuke Matsushita, Kengo Hirata, Ryo Wakizaka, and Emanuele D'Osualdo. 2026. RapunSL: Untangling Quantum Computing with Separation, Linear Combination and Mixing. *Proceedings of the ACM on Programming Languages* 10, POPL, Article 6 (2026). doi:10.1145/3776648

[13] Microsoft. 2025. Quantum Memory Management in Q#. Microsoft Learn, Azure Quantum documentation. https://learn.microsoft.com/en-us/azure/quantum/user-guide/language/statements/quantummemorymanagement

[14] Anouk Paradis, Benjamin Bichsel, Samuel Steffen, and Martin Vechev. 2021. Unqomp: Synthesizing Uncomputation in Quantum Circuits. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI '21)*. ACM, 222–236. doi:10.1145/3453483.3454040

[15] Anouk Paradis, Benjamin Bichsel, and Martin Vechev. 2024. Reqomp: Space-Constrained Uncomputation for Quantum Circuits. *Quantum* 8 (2024), 1258. doi:10.22331/q-2024-02-19-1258

[16] Robert Rand, Jennifer Paykin, Dong-Ho Lee, and Steve Zdancewic. 2019. ReQWIRE: Reasoning about Reversible Quantum Circuits. *Electronic Proceedings in Theoretical Computer Science* 287 (2019), 299–312. doi:10.4204/EPTCS.287.17

[17] P. L. Robinson. 2018. Julia Operators and Halmos Dilations. arXiv:1803.09329. doi:10.48550/arXiv.1803.09329

[18] Bonan Su, Li Zhou, Yuan Feng, and Mingsheng Ying. 2026. Borrowing Dirty Qubits in Quantum Programs. In *Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2 (ASPLOS '26)*. 274–289. doi:10.1145/3779212.3790134

[19] Bonan Su, Li Zhou, Yuan Feng, and Mingsheng Ying. 2026. An Interpretation of Bunched Logic for Reasoning about Heap-Manipulating Quantum Programs. arXiv:2409.10153v2 [cs.PL]

[20] John Watrous. 2018. *The Theory of Quantum Information.* Cambridge University Press. doi:10.1017/9781316848142

[21] Xiaoquan Xu, Chenke Liu, Boning Meng, Zihao Shen, and Li Zhou. 2026. Bona: Automatic Management of Dirty Ancilla Borrowing in Quantum Circuits. arXiv:2608.08765v1 [cs.PL]

[22] Charles Yuan, Christopher McNally, and Michael Carbin. 2022. Twist: Sound Reasoning for Purity and Entanglement in Quantum Programs. *Proceedings of the ACM on Programming Languages* 6, POPL, Article 30 (2022), 32 pages. doi:10.1145/3498691